

The background of the slide features a stylized, glowing blue wireframe profile of a human head facing right. Inside the head, there are intricate circuit-like patterns and binary code (0s and 1s) in various shades of blue and white, suggesting a digital or artificial intelligence theme. The overall color scheme is dark blue with bright blue highlights.

Machine Learning, Artificial Intelligence, and Big Data Analytics (IL, 4th Semester)

Lecture 1

Agenda (Lecture 1)

- About your teacher
 - Introduction round
 - Program of the Course
 - Organizational Information
-
- Intro to Data Science (AI, ML, and Big Data)
 - Supervised vs. Unsupervised Learning
 - The Data Science process
 - Recap of R

About me (Danilo)



- Born in Messina, Italy
- Studies
 - BSc Engineering at TU Bari
 - MSc Engineering at FH Technikum-Wien
 - PhD Computer Science at Uni Wien
- Work
 - Researcher @ Telecommunication Research Center Vienna (FTW)
 - Lecturer "Automotive Telecommunications" @ FH Technikum-Wien
 - Senior Research Scientist at Siemens AG Austria
- Interests
 - Technology, Innovation, and Trends; Data Science; AI ethics; Human Rights
- Private
 - 2 kids, play guitar, like biking/swimming, play videogames

Typical projects



Customer and Sales Analytics

Cross/up-selling, lead generation, pricing, customer behavior, ...



DSS for Product Life Cycle

Supply-chain diagnostic, product configuration, ...



Urban and Building analytics

Smart sensing, energy optimization, renewables, sustainability...



Industrial AI

Shopfloor monitoring, predictive maintenance, production optimization, ...

Research topics of interest

- Predictive Machine Learning in all its flavors (statistical learning, sub-symbolic learning, ...)
- Unsupervised learning, pattern recognition, anomaly detection, clustering, ...
- Recommender systems, collaborative filtering, content-based filtering, ...
- Visual analytics and UX, Explainable AI (XAI), Interpretable Machine Learning, ...
- (new) Neural-Symbolic AI.

About me (Stefan)



- Born in Vorarlberg
- Studies
 - BSc Software & Information Engineering @ TU Vienna
 - MSc Information & Knowledge Management @ TU Vienna
 - PhD Computer Science @ TU Vienna
- Work
 - Researcher @ DERI Galway, Ireland
 - Project Assistant @ TU Vienna and WU Vienna
 - Research Scientist @ Siemens AG Austria
- Interests
 - Knowledge Graphs, AI
- Private
 - 2 kids

Semantic Technologies and KGs

- Smart buildings
- Rail networks
- Statistical data
- IoT (meta) data
- Energy communities
- Company data
- Production data
- Product data

About you

- Anything you want to share about you, e.g.
 - Where are you from?
 - Why did you join this BSc program?
 - What are your interests?
 - Do you work in parallel to your studies?
 - ...

Setting the baseline



- Scan the barcode from your mobile phone

OR

- go to <http://sli.do> and insert this code: **30191**

and follow my instructions.

Program of the Course

- Recap and Introduction
- Module 1: Supervised Learning
- Module 2: Unsupervised Learning
- Module 3: Neural networks
- Practical hints, examples, and how to manage a ML project

- EXAM

Detailed Plan

Date	Lecture	hrs	Topic	Homework (tentative plan)
20.02.2023	Lecture 1	3	Overview AI/ML. Supervised vs. Unsupervised. The data science process. Recap of R	
27.02.2023	Lecture 2	3	Overfitting/Underfitting. Bias and Variance. Data Splitting. CrossValidation.	
06.03.2023	Lecture 3	3	Model Evaluation metrics. Intro to R/Python for data science.	
13.03.2023	Lecture 4	3	Classification: Decision Trees	Assignment Exercise 1
20.03.2023	Lecture 5	3	Classification: Ensemble learning. Random Forest, Adaboost, and xgboost	
27.03.2023	Lecture 6	3	Classification: K-Nearest Neighbors, Distance measures I	Assignment Exercise 2
20.04.2023	Lecture 7	3	Clustering: Intro and k-means (ATTENTION: Thursday)	Assignment Exercise 3
24.04.2023	Lecture 8	3	Clustering: Hierarchical clustering	
08.05.2023	Lecture 9	3	Clustering: Density-based clustering. Distance measures II.	Assignment Exercise 4
15.05.2023	Lecture 10	3	Intro to Neural Networks	
22.05.2023	Lecture 11	3	Intro to Deep Neural Networks	Assignment Exercise 5
05.06.2023	Lecture 12	3	ML in action. Wrap-up and preparation for the exam	
19.06.2023		2	EXAM	
	EL	2	Classification: SVM	
	EL	2	Hyperparameters tuning in R/Python	

Assumption (pre-requisite)




- Basic knowledge of statistics
 - Probability, conditional probability, Bayes, ...
 - Main distributions
 - Descriptive statistics (mean, median, mode, variance, ...)
- Basic programming skills
- Curiosity to try out new things and learn from data

Evaluation

- 40% Exercise
 - Homework
 - 5 exercises
 - Class work
 - After important topics (mostly small coding exercise)
 - Bonus given to those who present their solution in class
- 60% Written test

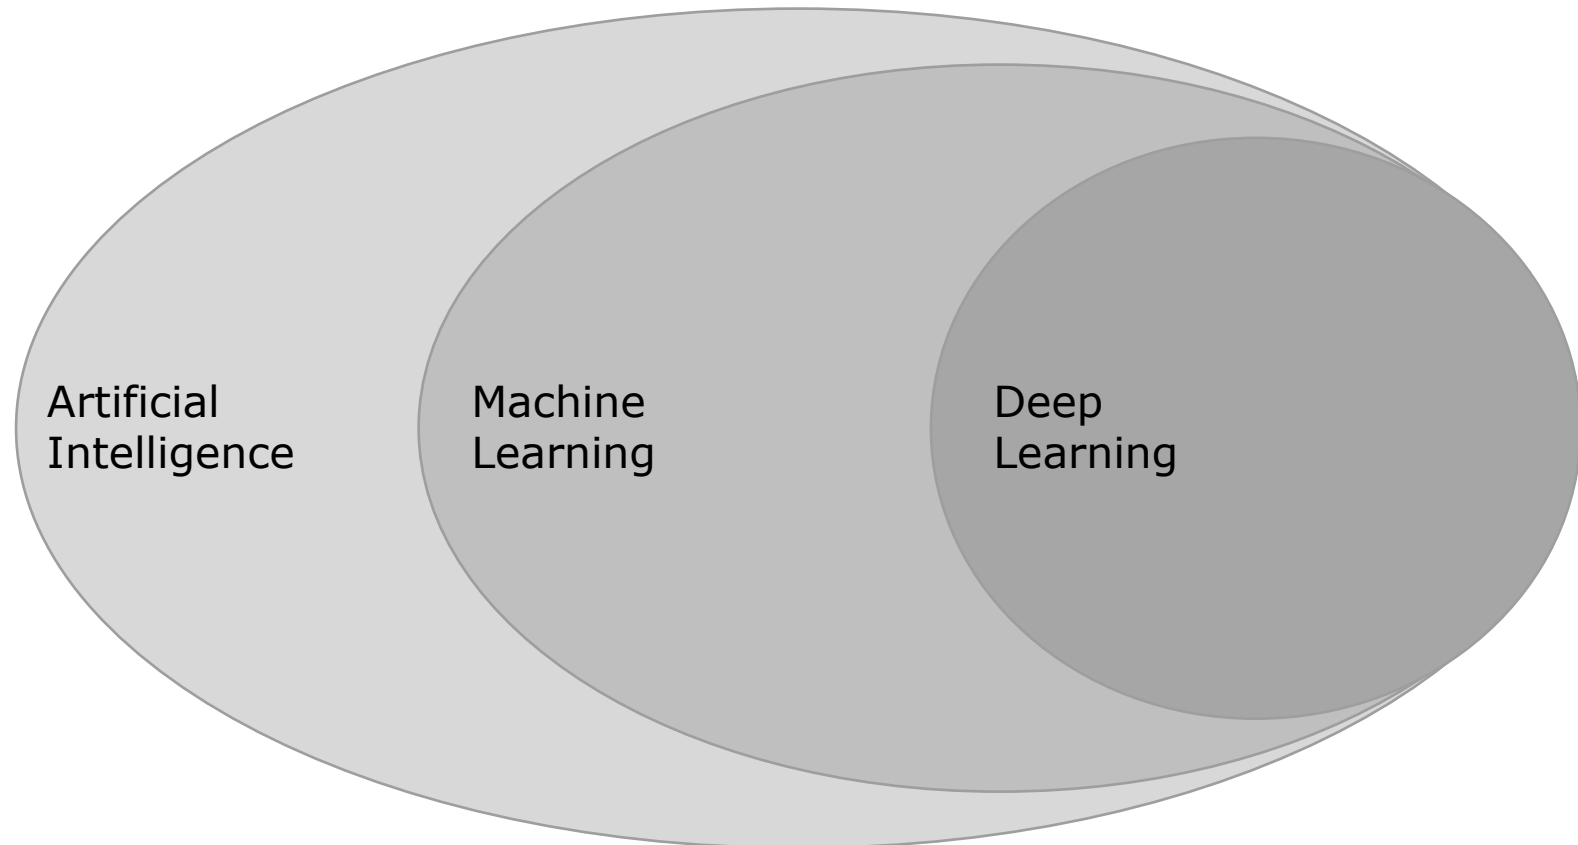
Introduction to data science

What is Artificial Intelligence




A term that has been misused too often...

What is Artificial Intelligence



Artificial Intelligence


Definition



- “The theory and development of computer systems **able to perform tasks normally requiring human intelligence**, such as visual perception, speech recognition, decision-making, and translation between languages.” (Oxford Living Dictionary)
- “The ability of a digital computer or computer-controlled robot to **perform tasks commonly associated with intelligent beings**.” (Encyclopedia Britannica)
- “The field of computer science dedicated to **solving cognitive problems commonly associated with human intelligence**, such as learning, problem solving, and pattern recognition.” (Amazon)

Artificial Intelligence

Types of AI



- **Symbolic AI** (aka classical AI or rule-based AI)

Relies on an explicit representation of a domain, hard-coded by humans as a set of symbols and rules. Uses deductive reasoning, logical inference, and other deterministic approaches to solve problems. Very useful in high-risk domains.

- **Non-Symbolic AI**

Tries to approach intelligence without specific representations of knowledge/domain. It learns autonomously by being *trained* with enough raw information to construct its own implicit knowledge. It can be further divided into sub-symbolic AI and statistical AI.


Artificial Intelligence

Types of AI (Example – TicTacToe)

X	O	O
O		X
X		

Artificial Intelligence


Types of AI (Example – TicTacToe)



- **Symbolic AI** (aka classical AI or rule-based AI)
 - Define and model the problem
 - Traverse possible solutions (e.g., search algorithm)
 - Deduct the best move
 - Challenge: Branching can make the possible solution space quickly increase requiring more sophisticated searches or modeling additional knowledge. (e.g. solution space for chess $\rightarrow 10^{120}$)

Artificial Intelligence


Types of AI (Example – TicTacToe)



- **Symbolic AI** (aka classical AI or rule-based AI)
 - Define and model the problem
 - Traverse possible solutions (search algorithm)
 - Deduct the best move
 - Challenge: Branching can make the possible solution space quickly increase requiring more sophisticated searches or modeling additional knowledge. (e.g. solution space for chess $\rightarrow 10^{120}$)
- **Non-Symbolic AI**
 - Observe previously played games
 - Infer the move that leads to higher chance of success.
 - Challenge: The data about previously played games could not be sufficient to make an inference

Machine learning

Definition




- “the use and development of computer systems that are able **to learn and adapt without following explicit instructions**, by using algorithms and statistical models to analyze and draw inferences from patterns in data” (Oxford living Dictionary)
- “Machine learning (in AI) is a discipline concerned with the implementation of computer software that **can learn autonomously**” (Encyclopedia Britannica)
- “A subset of artificial intelligence (AI) that provides systems the ability to **automatically learn and improve from experience** without being explicitly programmed.” (IBM)

Machine learning

What is "learning"


"A computer program is said to learn from experience ***E*** with respect to some class of tasks ***T*** and performance measure ***P*** if its performance at tasks in ***T***, as measured by ***P***, improves with experience ***E***."

Tom Mitchell (*Machine Learning*. McGraw Hill).



Machine learning

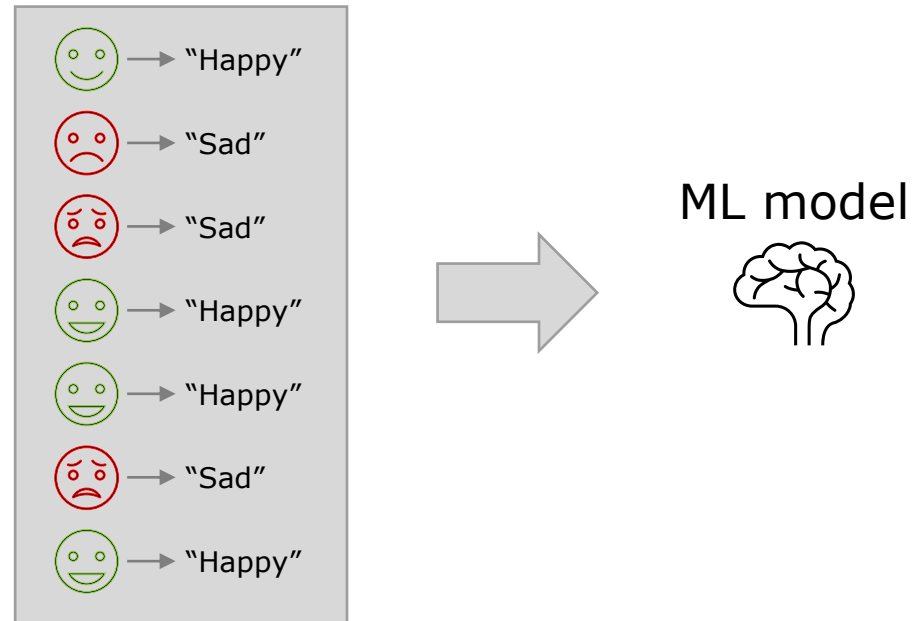
Types of ML



- **Supervised Learning:** The machine is presented with a set of input data and output data and learns the relation between input and output such that it can predict the output for new input data.
- **Unsupervised Learning:** The machine is presented with just input data and autonomously search for structure and patterns within the input data.
- **Semi-supervised Learning:** A mix of the above.
- **Reinforcement Learning:** The machine does not only learn from a static dataset but continuously learns through trial after receiving positive or negative feedback as reinforcement.

Machine learning

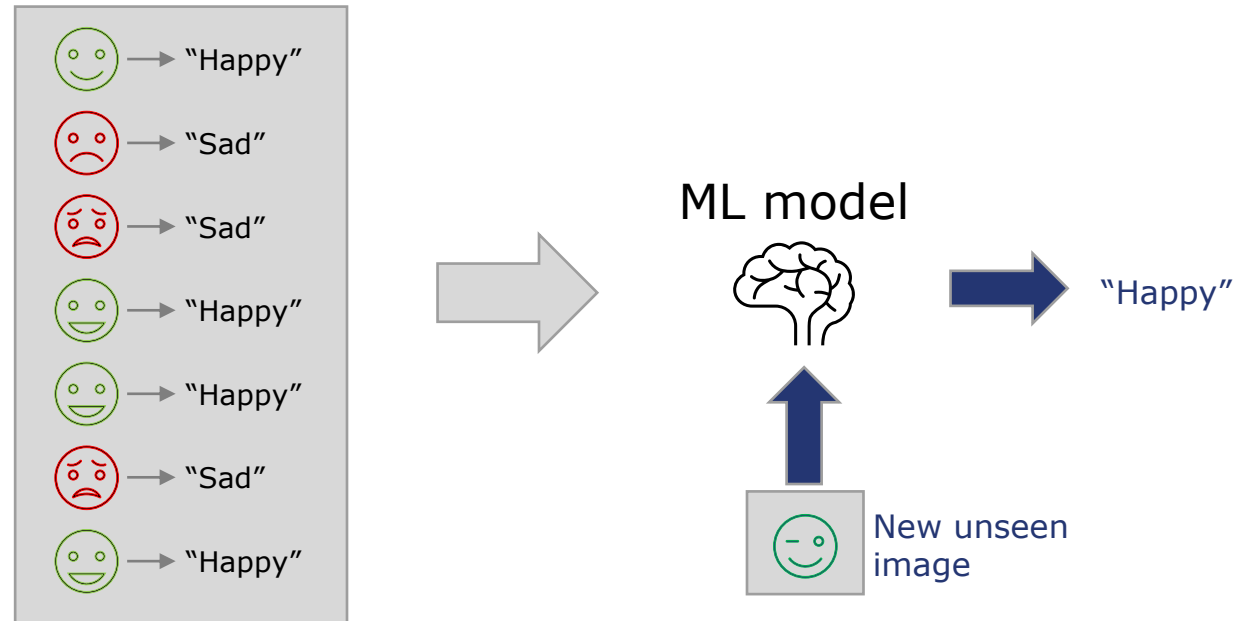
Supervised learning



Training data (contains both input and output)

Machine learning

Supervised learning

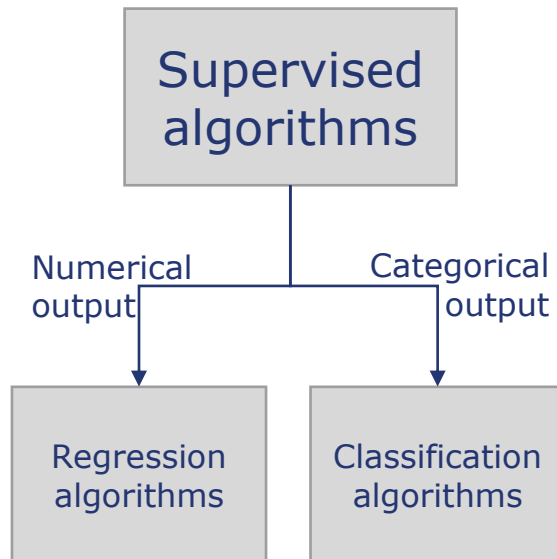


Training data (contains both input and output)

Machine learning

Supervised learning

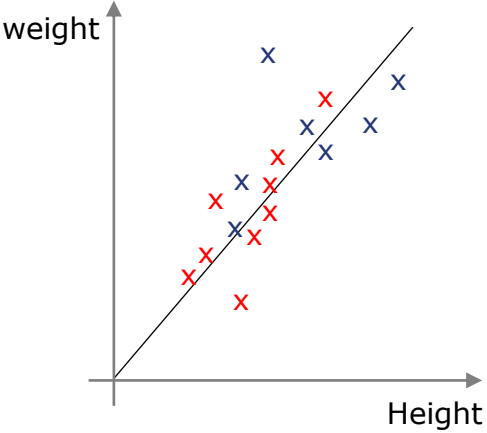
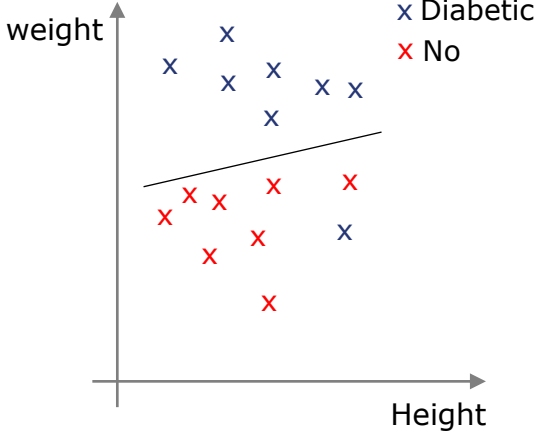
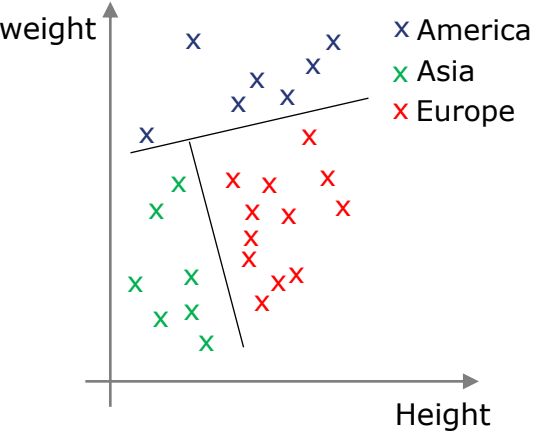
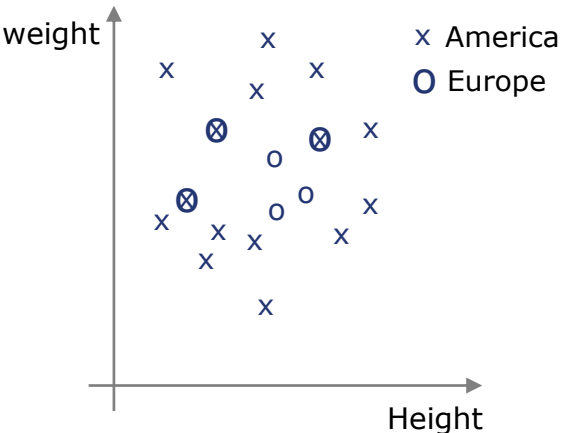
Two main types of supervised learning algorithms



- **Regression:** The variable to predict is numerical
 - Examples: Predict the price of an object. Predict the age of a subject. Predict the income of a person. Predict the cost of a project. Predict the lifetime of a device. Etc.
- **Classification:** The variable to predict is categorical (a label)
 - Examples: Spam detection ("spam" vs. "not spam"), image recognition ("table" vs. "chair" vs. "TV" vs. ...), malfunctioning component ("proper" vs. "malfunctioning"), credit reliability of a person ("reliable" vs "unreliable"), etc.


Machine learning

Supervised learning examples


Regression	Classification		
	Binary	Multi-class	Multi-label
<p>Predict weight based on height</p> 	<p>Predict if subject has diabetes or not</p> 	<p>Predict in which continent the subject is living</p> 	<p>Predict in which continent subject has lived</p> 

Machine learning



Unsupervised learning



Training data (contains only input)

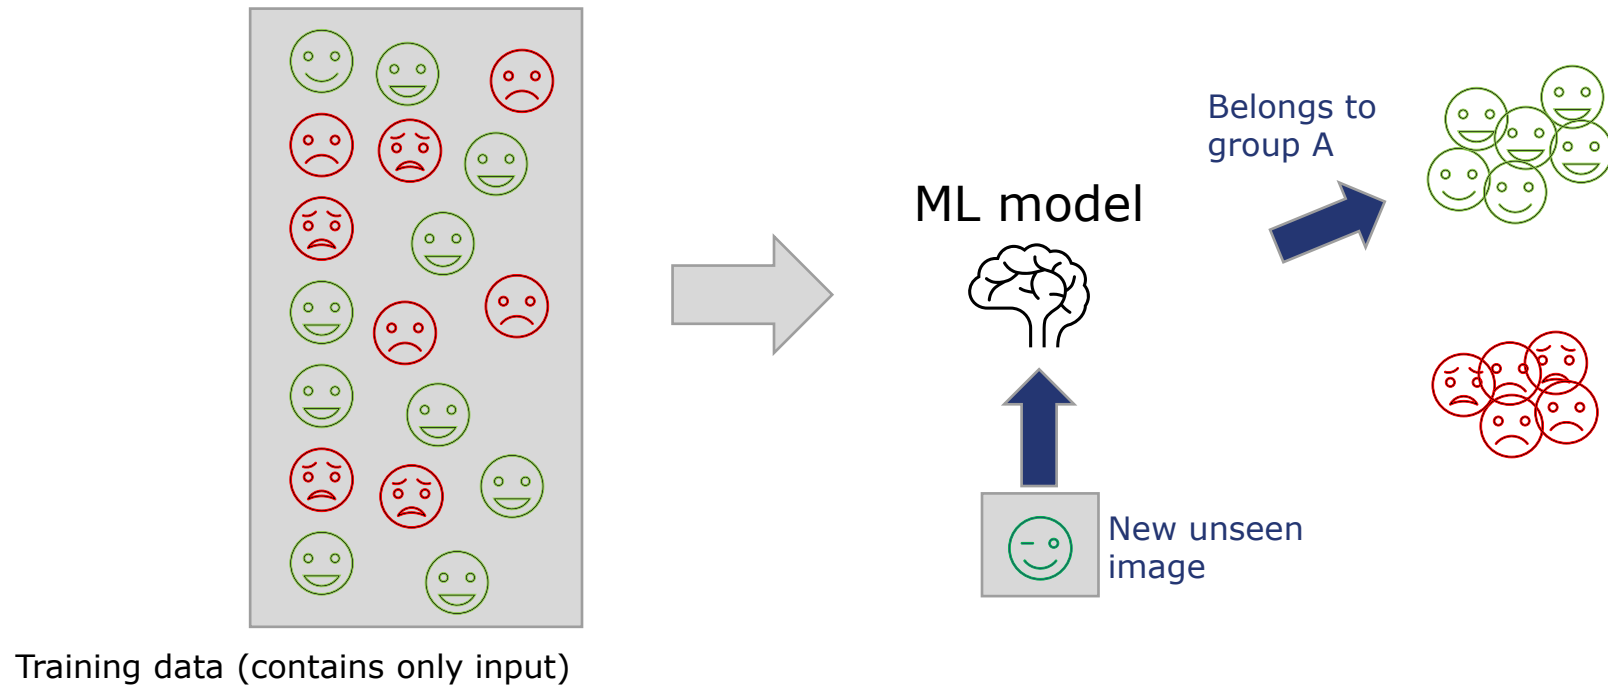


ML model



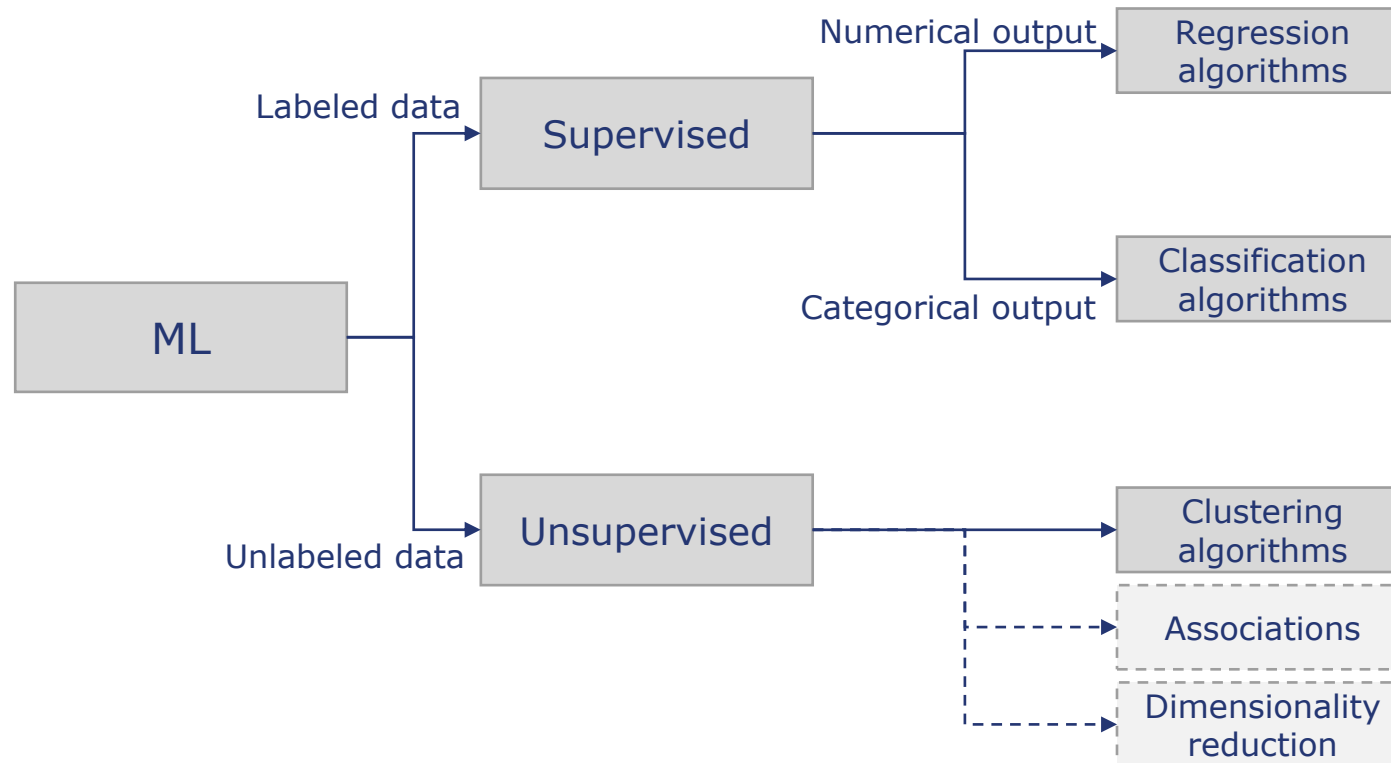
Machine learning

Unsupervised learning



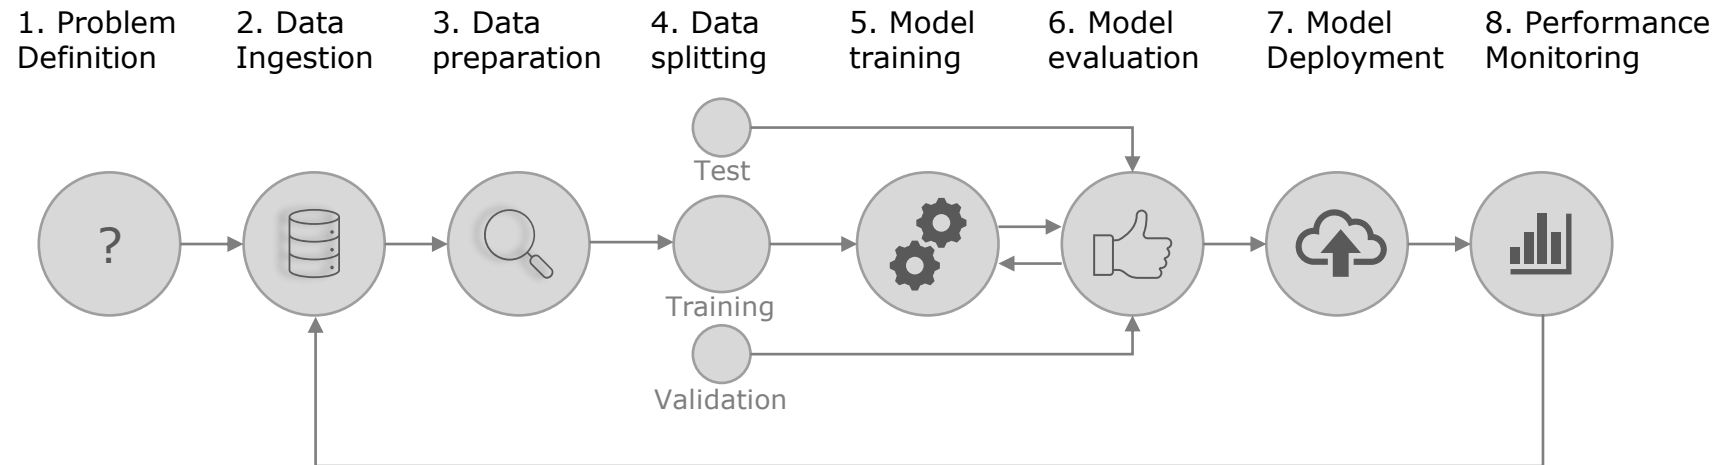
Machine Learning

Supervised vs. Unsupervised Summary

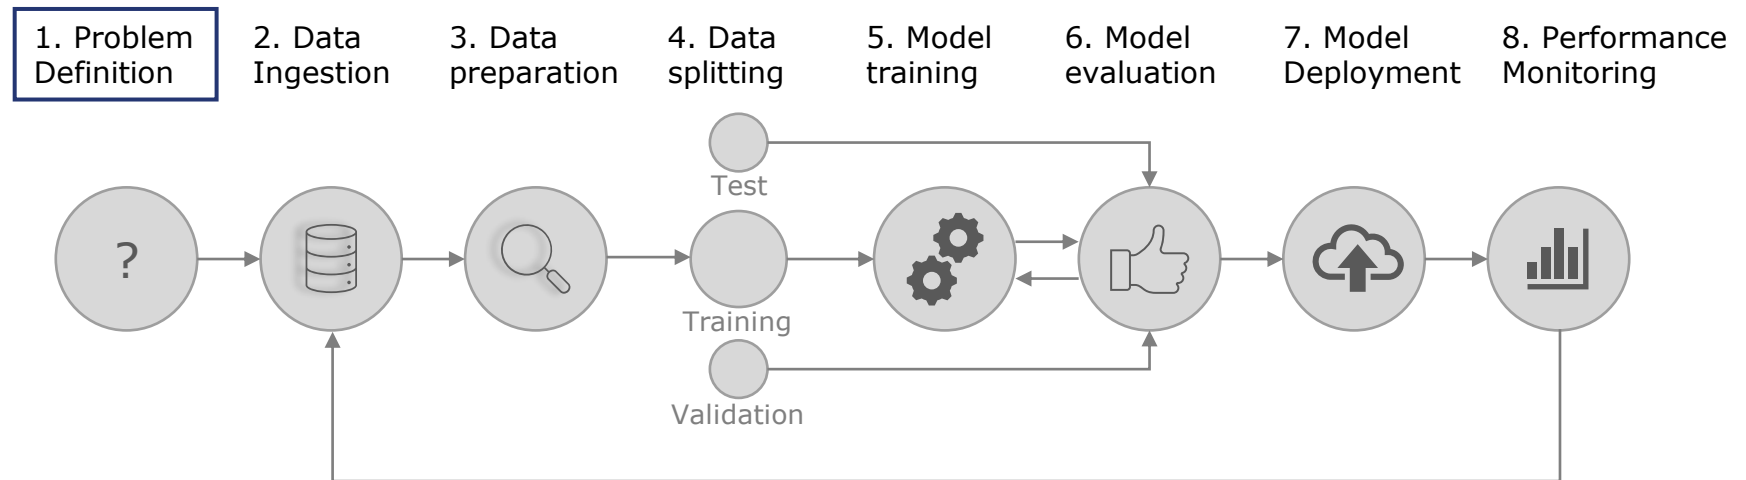


- Linear regression
 - Simple
 - Multiple
 - Multivariate
- Non-linear regression
- Decision Tree
- Random Forest
- Naïve Bayes
- Logistic regression
- Support Vector Machines
- K-nearest neighbors
- Decision Tree
- Random Forest
- ...
- K-means
- Hierarchical clustering
- Density-Based clustering
- Model-based clustering
- ...

The Machine Learning pipeline

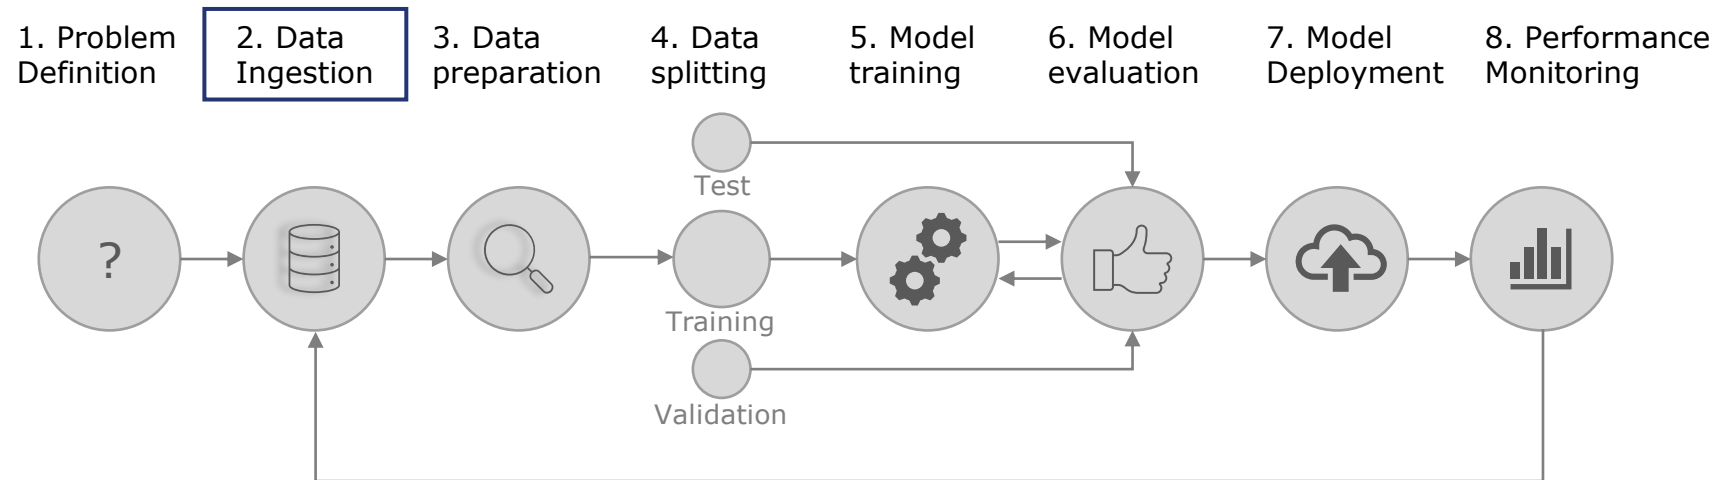


The Machine Learning pipeline



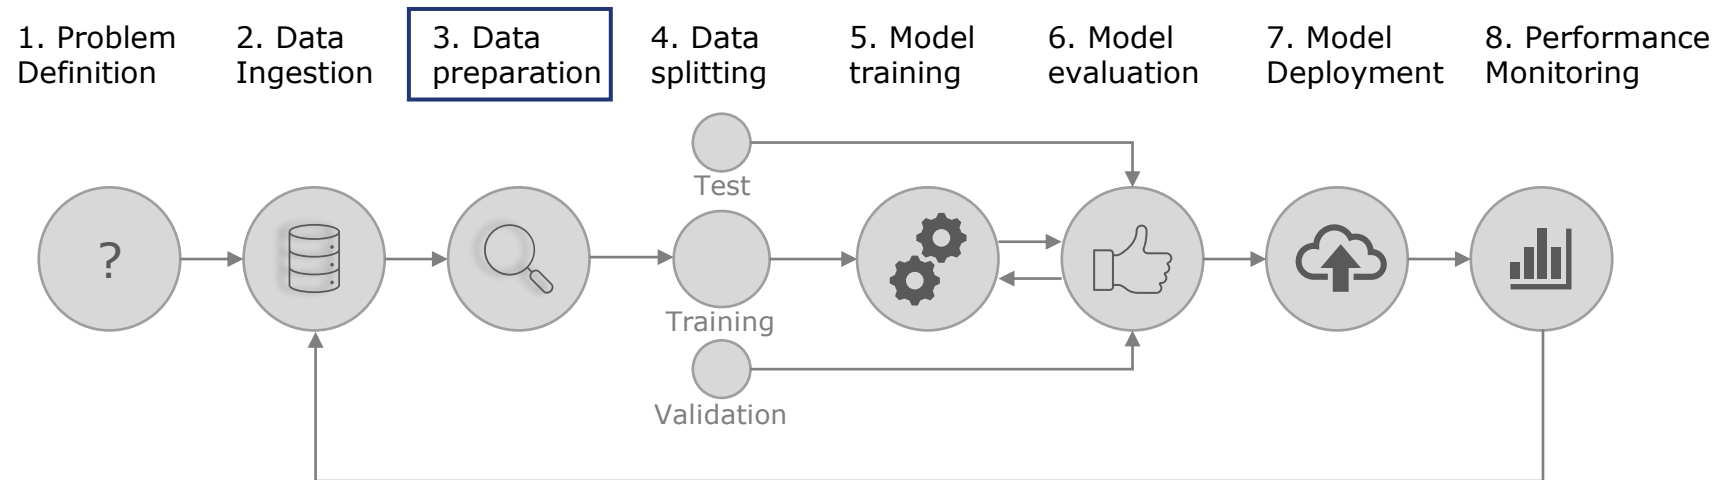
- Define the Business Question.
- Learn the domain.
- Translate the business question into a data science question.

The Machine Learning pipeline



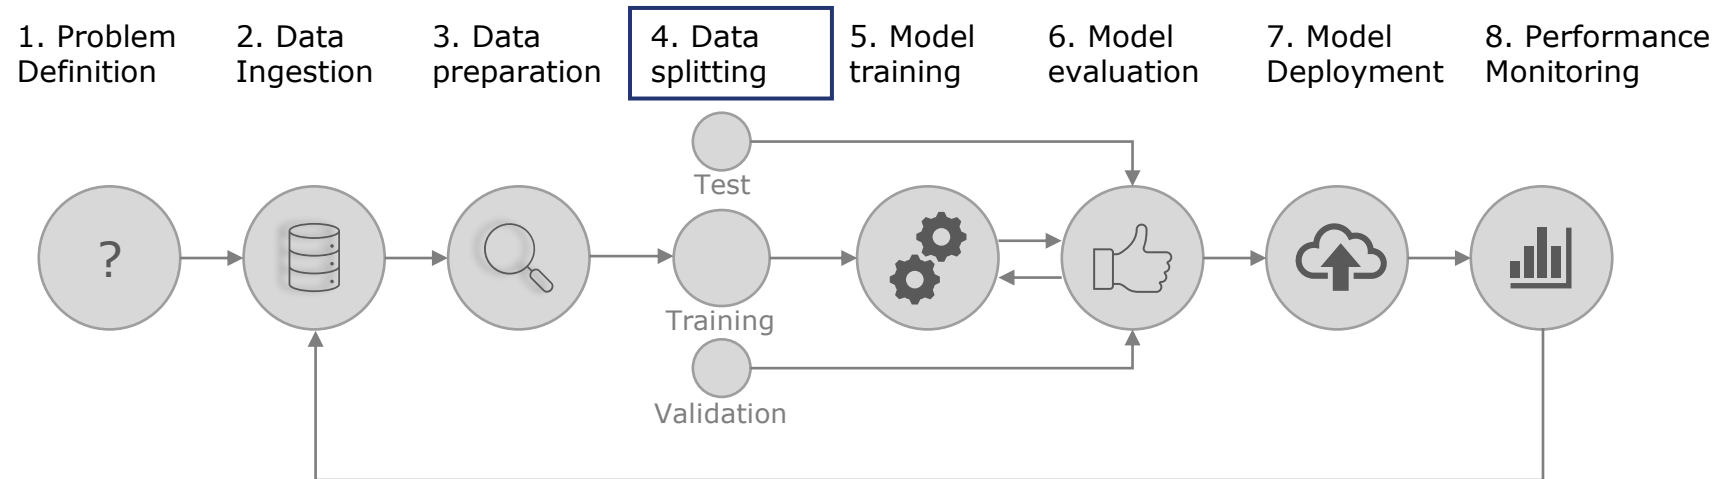
- Identify which data is necessary and how to collect it
- Offline vs. Online
- Database vs. Files

The Machine Learning pipeline



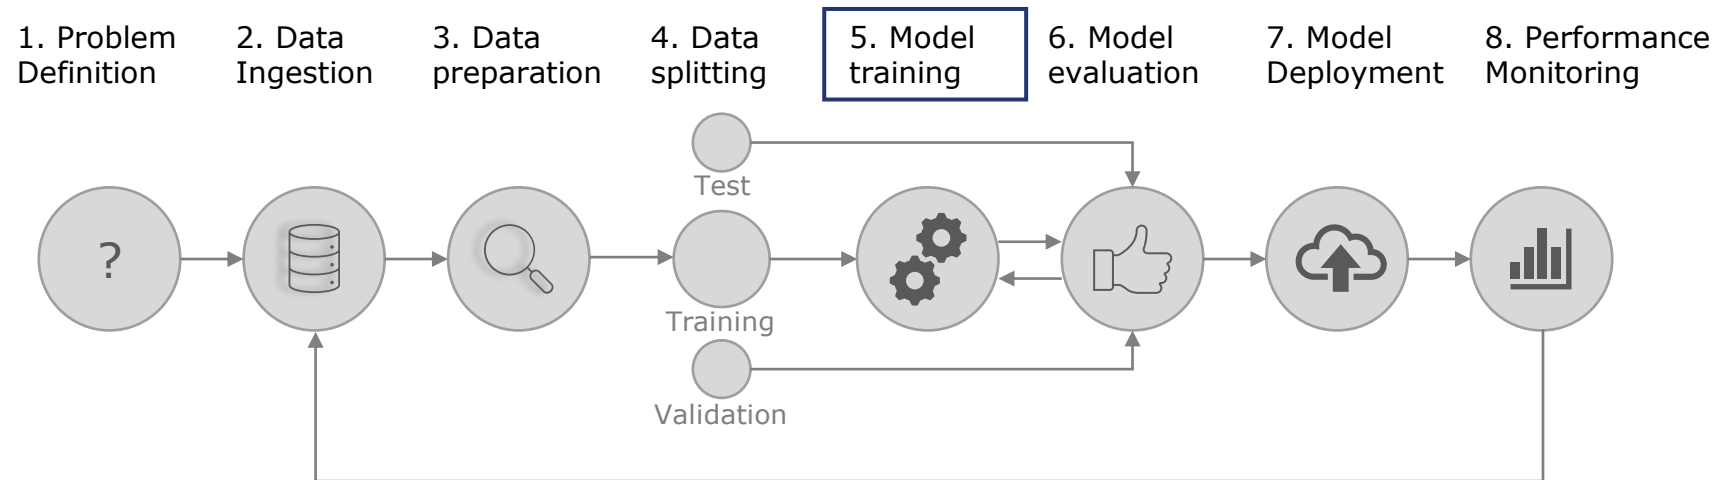
- Cleansing (remove erroneous entries, fill missing values, remove duplicates, etc.)
- Feature selection
- Feature engineering

The Machine Learning pipeline



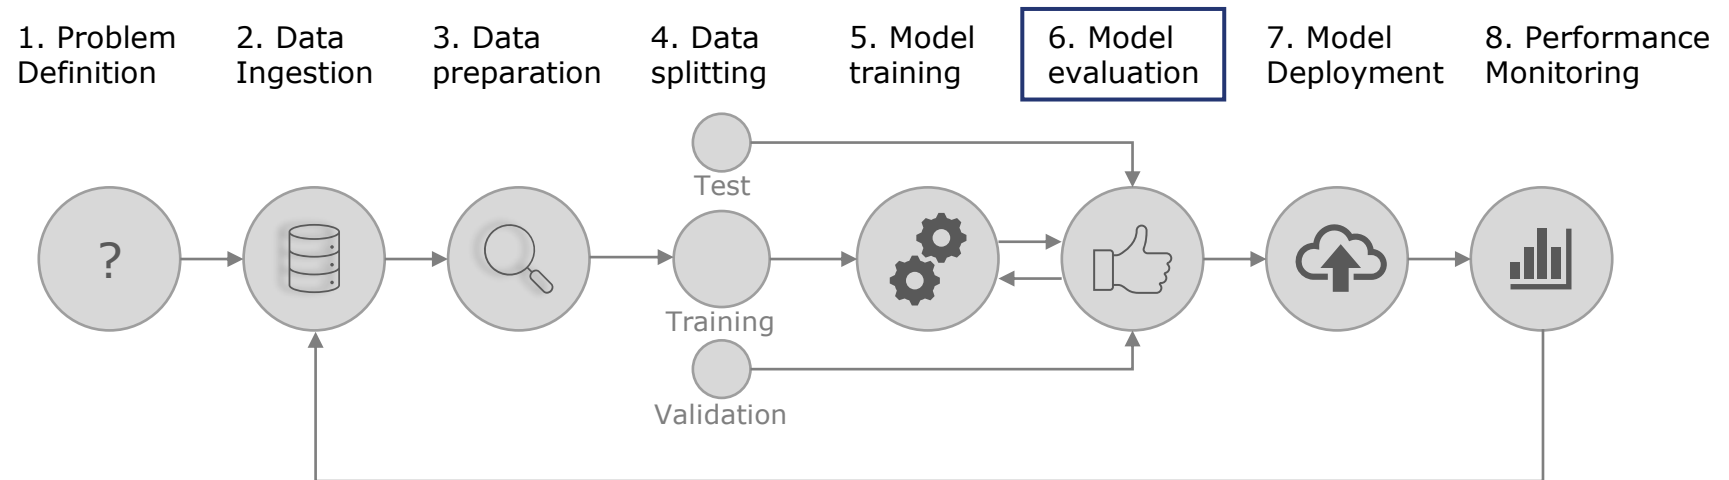
- Split into two or three subsets
- *Training set* → Used to train your model
- (*Validation set*) → Used for validate the trained model (e.g., tuning the parameters)
- *Test set* → Used to the evaluate the final tuned model to other final models

The Machine Learning pipeline



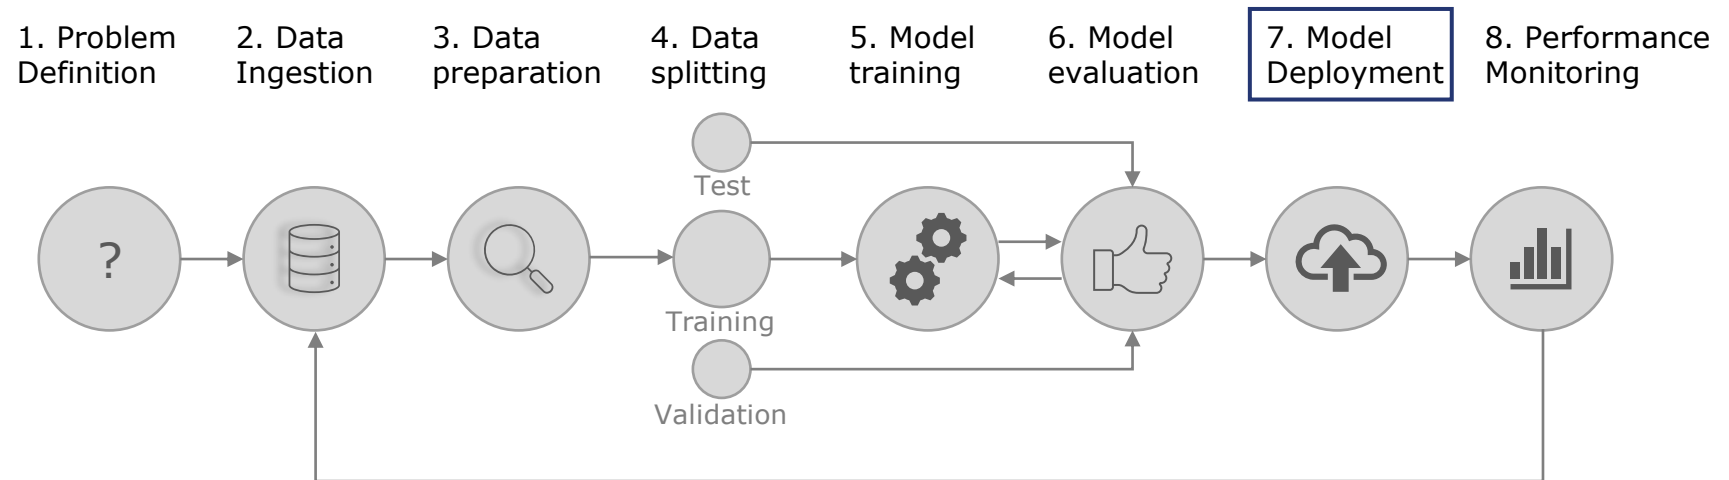
- Identify which type of algorithm is appropriate for training your model
 - Desired Accuracy,
 - Desired Interpretability,
 - Desired Scalability,
 - Constraints on processing power

The Machine Learning pipeline



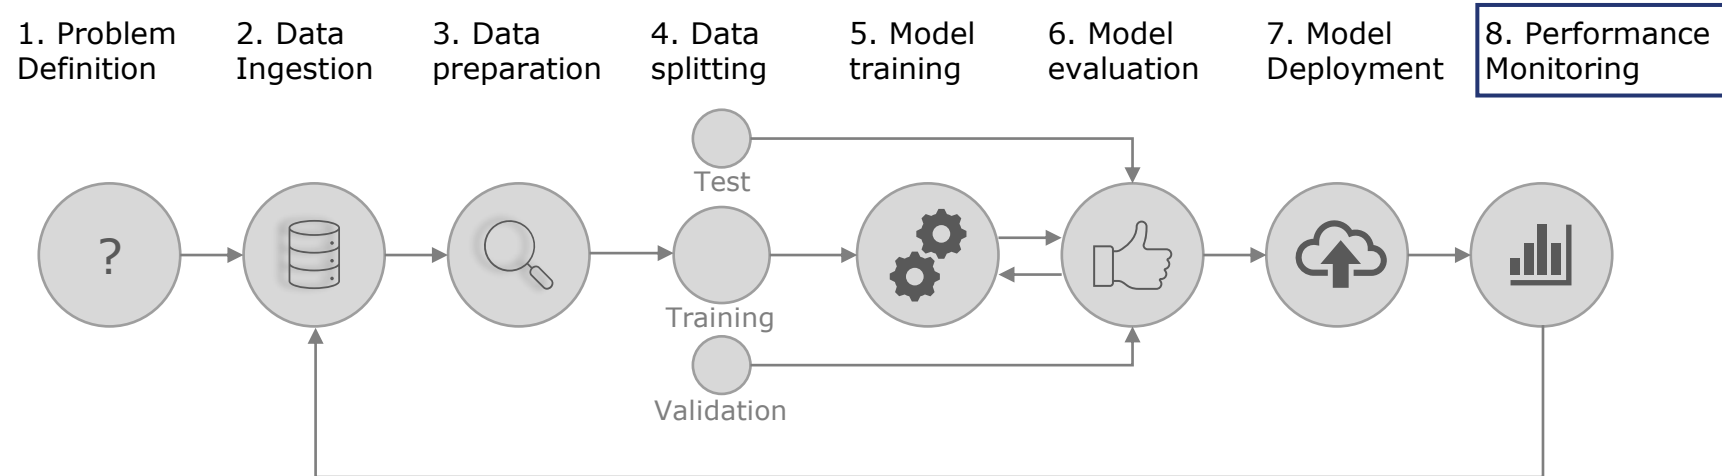
- Use the validation and test data sets to assess the model accuracy
- Iterate 5. and 6. until appropriate

The Machine Learning pipeline



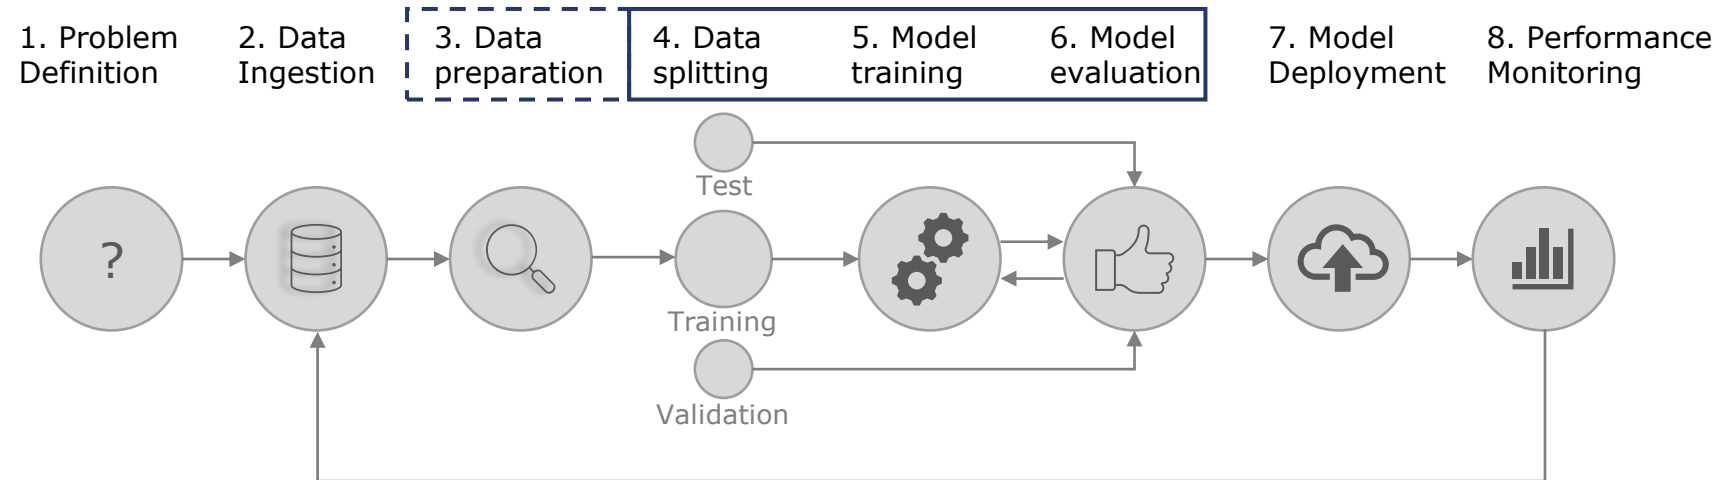
- Define how to deploy the model. API? Dashboard? Integrated Pipeline? Etc.

The Machine Learning pipeline



- Continuously monitor the performance of the model to identify deviations and deterioration of accuracy.

Focus of this course



Recap of R

- Setup for the course
- The R language
 - *R* ($\geq 4.0.0$, suggested $\geq 4.2.x$) – <https://cran.r-project.org/>
- Common IDEs (pick one)
 - *RStudio* (Simply the best. I strongly suggest to update to v2022.xx.x)
 - *Visual Studio Code* (the emerging one. Requires [this](#))
 - *RGui* (the lightest. Comes with R. Lacks several functionalities)
 - *Jupyter Notebook* or *Jupyter Lab*
- During the course we'll also use **Python**... and if we have time we'll introduce **Julia** as well.

Recap of R

Lecture continues on Rstudio